



Getting Data Into Snowflake

Snowflake Activation Webinar Series

Snowflake Account Engineering Team

Safe Harbor and Disclaimers

Other than statements of historical fact, all statements contained in these materials and any accompanying oral commentary (collectively, the “Materials”) are forward-looking statements within the meaning of Section 27A of the Securities Act of 1933, as amended, and Section 21E of the Securities Exchange Act of 1934, as amended, including statements regarding (i) Snowflake’s business strategy, plans, opportunities, or priorities; (ii) Snowflake’s new or enhanced products, services, and technology offerings, including those that are under development or not generally available; (iii) market size and growth, trends, and competitive considerations; (iv) Snowflake’s vision, strategy, and expected benefits relating to artificial intelligence (“AI”), Snowpark, Snowflake Marketplace, the AI Data Cloud, and AI Data Clouds for specific industries or product categories, including the expected benefits and network effects of the AI Data Cloud; and (v) the integration, interoperability, and availability of Snowflake’s products, services, or technology offerings with or on third-party platforms or products, including public cloud platforms. These forward-looking statements are subject to a number of known and unknown risks, uncertainties, and assumptions, including those described under the heading “Risk Factors” and elsewhere in the Annual Reports on Form 10-K and the Quarterly Reports on Form 10-Q that Snowflake files with the Securities and Exchange Commission from time to time. In light of these risks, uncertainties, and assumptions, the future events and trends discussed in the Materials may not occur, and actual results could differ materially and adversely from those expressed or implied in the forward-looking statements. As a result, you should not rely on any forward-looking statements as predictions of future events. Forward-looking statements speak only as of the date the statements are first made and are based on information available to Snowflake at the time those statements are made and/or management’s good faith belief as of that time. Except as required by law, we undertake no obligation, and do not intend, to update the forward-looking statements in these Materials.

Any future product or roadmap information (collectively, the “Roadmap”) is intended to outline general product direction. The Roadmap is not a commitment, promise, or legal obligation for Snowflake to deliver any future products, features, or functionality; and is not intended to be, and shall not be deemed to be, incorporated into any contract. The actual timing of any product, feature, or functionality that is ultimately made available may be different from what is presented in the Roadmap. The Roadmap information should not be used when making a purchasing decision. In case of conflict between the information contained in the Materials and official Snowflake documentation, official Snowflake documentation should take precedence over these Materials. Further, note that Snowflake has made no determination as to whether separate fees will be charged for any future products, features, and/or functionality which may ultimately be made available. Snowflake may, in its own discretion, choose to charge separate fees for the delivery of any future products, features, and/or functionality which are ultimately made available.

The Materials may contain information provided by third-parties. Snowflake has not independently verified this information, and usage of this information does not mean or imply that Snowflake has adopted this information as its own or independently verified its accuracy.



YOUR SESSION ROADMAP

- Choosing a Load Path
- Staging & File Formats
- Batch Load with COPY INTO
- Continuous Load: Snowpipe
- CoCo-Assisted Migration



DATA ONLY WORKS ONCE IT'S IN


Match the path to the source before writing pipeline code.

- A one-off load script breaks when volume, arrival pattern, or schema changes.
- The cost appears as rework, stale tables, and blocked downstream users.
- This lab chooses a path, inspects files, loads data, and handles a new column.



THE INGESTION LANDSCAPE


Start with what the source produces, then choose the matching path.



SOURCE SHAPE → INGESTION PATH → LOADED TABLE

BULK OR CONTINUOUS? FOLLOW THE FILE

Both paths load files. Arrival pattern determines who initiates the load.



MATCH THE PATH TO THE SOURCE

COPY INTO

TRIGGER
Your schedule

MANAGED BY
Your orchestration

BEST FOR
Periodic bulk loads

SNOWPIPE

TRIGGER
Cloud event

MANAGED BY
Snowflake

BEST FOR
Continuous file arrival

OPENFLOW OR PARTNER

TRIGGER
Connector sync

MANAGED BY
Connector platform

BEST FOR
Database, SaaS, or custom sources



SOURCE-FIRST DECISION TREE

Choose the ingestion path by asking what the source produces.

1 Does the source produce files?

→ **COPY INTO or Snowpipe**

2 Does it emit rows or events?

→ **Snowpipe Streaming**

3 Is it a database or SaaS app?

→ **Openflow or partner connector**

4 Is it a legacy warehouse migration?

→ **CoCo-assisted migration**

LIVE DEMO PATH • FILES → STAGE → COPY INTO



OPENFLOW OPERATING MODEL

A connector moves data; a runtime operates the flow.

SOURCE SYSTEMS

DATABASES

SaaS + APIs

FILES + EVENTS



OPENFLOW
CONNECTOR
Source-specific movement

+

RUNTIME
Operate the flow

MANAGED • GOVERNED • OBSERVABLE



SNOWFLAKE

TABLES


ANALYTICS

AI

PARTNER PATH • Use an approved connector when it best fits the source and operating model.

STAGE + FILE FORMAT = REUSABLE LOAD CONTRACT

Separate location from parsing rules, then reuse both across loads.



COPY INTO reuses the named stage and file format.



CHOOSE THE FILE FORMAT

The format controls how Snowflake interprets fields, nesting, and logical types.

CSV

Flat exports
Broad tool support

Watch:
headers, quotes,
nulls

JSON

APIs and nested data
Flexible structure

Watch:
arrays and paths


PARQUET

Columnar analytics
Logical types

Watch:
schema consistency

INSPECT → INFER → CREATE

Let staged files inform the table definition before the first load.



BATCH LOAD WITH COPY INTO

Inspect → infer → create → load → evolve.

1
INSPECT

2
INFER

3
CREATE

4
LOAD

5
EVOLVE

CONTROL THE SCHEMA BEFORE AND AFTER THE FIRST LOAD



SCHEMA EVOLUTION: BEFORE AND AFTER

A new source column becomes a new table column while earlier rows remain intact.




TABLE: ENABLE_SCHEMA_EVOLUTION = TRUE • LOAD: MATCH_BY_COLUMN_NAME
ROLE: EVOLVE SCHEMA or OWNERSHIP
CSV: PARSE_HEADER • ERROR_ON_COLUMN_COUNT_MISMATCH = FALSE

NEW FILES ARE NOT ROW-LEVEL CHANGES

Use file ingestion for new files. Use CDC when the source changes individual records.

FILE INGESTION

New file arrives



COPY INTO / Snowpipe



Rows appended to table

CHANGE DATA CAPTURE

INSERT / UPDATE / DELETE



CDC-capable connector




Record changes replicated

Editing an already processed file does not
become row-level change capture.

NEW FILE → EVENT → SNOWPIPE → TABLE

Cloud events notify Snowpipe; the notification identifies the file, not the data.




FILE PATH: event-driven • ROW PATH: Snowpipe Streaming



SNOWPIPE: SETUP AND RECOVERY

Configure the event path once. Use refresh only for bounded recovery.



RECOVER MISSED FILES
ALTER PIPE ... REFRESH for bounded recovery, then monitor load history.

COCO GUIDES THE MIGRATION

Six documented stages keep conversion, assessment, and deployment explicit.




SNOWCONVERT = DETERMINISTIC CODE CONVERSION
CoCo guides assessment and remediation around that conversion layer.



FROM SOURCE TO SNOWFLAKE

Choose the path, establish the contract, then monitor the result.



THANK YOU



© 2026 Snowflake Inc. All Rights Reserved